

Lecture 2: Regression

Mark Hendricks

Summer 2023

FINM 25000

Outline

Regression

OLS Mathematics



THE UNIVERSITY OF
CHICAGO

Regression analysis in finance

Regression applications in finance include...

- ▶ **Risk-management.** Find how a portfolio return is impacted by some factor/instrument.
- ▶ **Forecasting.** Build forecasts of financial and macroeconomic variables. (inflation, yields, etc.)
- ▶ **Pricing.** The fundamental asset pricing equation is a linear relation between risk and return.



Beyond regression

Nonlinear analysis is also important.

- ▶ **Options Pricing.** Differential equations requiring martingale methods, simulation, finite differenc, etc.
- ▶ **Value at Risk.** Model the tail of the distribution of profits and losses.
- ▶ **Volatility Models** Need non-linear timeseries models such as GARCH.



Linear regression model

Consider a **linear regression model** involving two variables, y and x .

$$y = \alpha + \beta x + \epsilon$$

- ▶ y is referred to as the **regressand**, or explained variable.
- ▶ x is referred to as the **regressor**, covariate, or explanatory variable.
- ▶ α and β are the (constant) parameters of the model.



Example: Portfolio factor sensitivity

Decompose the hedge fund return into a market-driven and market-neutral return.

$$r_p = \alpha + \beta r_{\text{mkt}} + \epsilon$$

- ▶ random total portfolio return denoted by r_p
- ▶ random return on the S&P 500, denoted by r_{mkt} .

Interpret...

- ▶ $\beta = 0, 1, 2$
- ▶ $\alpha = -.01, 0, .01$.



Example: Portfolio decomposition

Continuing the example from above,

$$r_p = \alpha + \beta r_{\text{mkt}} + \epsilon$$

We may want to know “how much” of r_p is explained by r_{mkt} .

- ▶ **R-squared** (R^2) is a metric of the variation explained in the regression model.
- ▶ Is the hedge-fund driven by market returns if $\beta = 1$, $R^2 = .10$? How about $\beta = .5$, $R^2 = .50$?

(Notation: R^2 is standard notation in regression analysis—nothing to do with my choice of variable name r_p, r_{mkt} .)



Univariate regression

When there is only one regressor, x , we will see that the OLS estimator is simply:

$$\beta = \frac{\text{cov}(y, x)}{\text{var}(x)}$$

And that the R-squared statistic is simply

$$R^2 = [\text{corr}(y, x)]^2$$

So why bother with regression if we just need covariances and variances?



Multiple regression

In the case of multiple regressors, the OLS statistics are not so easily formed.

- ▶ Augment our hedge-fund regression with a second regressor: a US dollar index, $r_{\$}$.

$$r_p = \alpha + \beta_1 r_{\text{mkt}} + \beta_2 r_{\$} + \epsilon$$

- ▶ The formulas for β_1 and β_2 do not follow as easily:

$$\beta_1 \neq \frac{\text{cov}(r_p, r_{\text{mkt}})}{\text{var}(r_{\text{mkt}})}$$

- ▶ The R-squared stat captures the correlation between r_p and the combined space spanned by both r_{mkt} and $r_{\$}$.



Caution!

Remember that the multi-variable beta is not the same as the univariate beta!

$$r_p = \alpha + \beta_1 r_{\text{mkt}} + \beta_2 r_{\$} + \epsilon$$

- ▶ Perhaps r_p is positively correlated with $r_{\$}$, and thus would have a positive beta if regressed on only $r_{\$}$.
- ▶ But β_2 is not a measure of this pairwise comovement!
- ▶ β_2 gives the impact on r_p if we hold r_{mkt} constant!
- ▶ Thus, when the regressors are correlated, multi-variable betas can be quite different from their univariate counterpart.



Units

When interpreting the regression coefficients, be careful to remember the underlying units.

$$r_p = \alpha + \beta_1 r_{\text{mkt}} + \beta_2 r_{\$} + \epsilon$$

- ▶ The volatility of r_{mkt} is three times larger than the volatility of $r_{\$}$.
- ▶ Thus, even if β_2 is larger than β_1 , we need to remember that one-unit changes in $r_{\$}$ happen less frequently.
- ▶ In this situation it may be more helpful to report $\beta_1 \sigma_1$ and $\beta_2 \sigma_2$ to help convey the one-standard deviation impact from each factor.



Outline

Regression

OLS Mathematics



THE UNIVERSITY OF
CHICAGO

Multivariate linear regression

In a multivariate regression model with k regressors,

$$y = \alpha + \beta_1 x_1 + \beta_2 x_2 + \dots \beta_k x_k + \epsilon$$

$$= \alpha + \sum_{j=1}^k \beta_j x_j + \epsilon$$

$$= \mathbf{x}'\boldsymbol{\beta} + \epsilon$$

- ▶ The last line defines \mathbf{x} such that the first element is the constant 1, and the first element of $\boldsymbol{\beta}$ is α .
- ▶ Including the regression constant in the vector notation will simplify the algebra, as we will always consider the case where the first regressor is a constant.



THE UNIVERSITY OF
CHICAGO

Data from the regression model

A sample of n observations is denoted as (y_i, \mathbf{x}_i) for $i = 1, 2, \dots, n$.

$$y_i = \mathbf{x}_i' \boldsymbol{\beta} + \epsilon_i$$

where

$$\mathbf{x}_i \equiv \begin{bmatrix} 1 \\ x_{i,1} \\ x_{i,2} \\ \vdots \\ x_{i,k} \end{bmatrix} \quad \boldsymbol{\beta} \equiv \begin{bmatrix} \alpha \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_k \end{bmatrix}$$



Regression estimate

Consider a sample estimate of β , denoted by \mathbf{b} .

Then

$$y_i = \mathbf{x}_i' \mathbf{b} + e_i$$

where e_i denotes a sample residual,

$$e_i = y_i - \mathbf{x}_i' \mathbf{b}$$

This is estimated regression, as opposed to the population regression equation above.



Ordinary least squares

The **ordinary least squares estimator** of β minimizes the sum of squared sample errors:

$$\begin{aligned}\mathbf{b} &\equiv \arg \min_{\mathbf{b}_o} \sum_{i=1}^n (e_i)^2 \\ &= \arg \min_{\mathbf{b}_o} \sum_{i=1}^n (y_i - \mathbf{x}_i' \mathbf{b}_o)^2\end{aligned}$$



OLS problem

Rewrite the OLS problem in matrix notation,

$$\begin{aligned}\mathbf{b} &\equiv \arg \min_{\mathbf{b}_o} \mathbf{e}'\mathbf{e} \\ &= \arg \min_{\mathbf{b}_o} (\mathbf{Y} - \mathbf{X}\mathbf{b}_o)'(\mathbf{Y} - \mathbf{X}\mathbf{b}_o)\end{aligned}$$

where

$$\mathbf{X} \equiv \begin{bmatrix} \mathbf{x}'_1 \\ \mathbf{x}'_2 \\ \vdots \\ \mathbf{x}'_n \end{bmatrix}, \quad \mathbf{Y} \equiv \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}, \quad \mathbf{e} \equiv \begin{bmatrix} e_1 \\ e_2 \\ \vdots \\ e_n \end{bmatrix},$$



Assumption: Full-rank

Assumption 1: $\mathbf{X}'\mathbf{X}$ is full rank.

Equivalently, assume that there is no exact linear relationship among any of the regressors.

- ▶ Clearly, the existence of OLS estimator requires that this assumption be satisfied.
- ▶ Multicollinearity refers to the case where this assumption fails.



OLS estimate

Solving the minimization problem above gives the **OLS estimate**:

$$\mathbf{b} = (\mathbf{X}'\mathbf{X})^{-1} \mathbf{X}'\mathbf{Y}$$

- ▶ This estimate yields sample residuals of

$$\begin{aligned}\mathbf{e} &= \mathbf{Y} - \mathbf{X} (\mathbf{X}'\mathbf{X})^{-1} \mathbf{X}'\mathbf{Y} \\ &= \left(\mathbf{I} - \mathbf{X} (\mathbf{X}'\mathbf{X})^{-1} \mathbf{X}' \right) \mathbf{Y}\end{aligned}$$

- ▶ Thus \mathbf{e} is orthogonal to \mathbf{X} .
- ▶ Equivalently, the in-sample correlation between x_i and e_i is zero.



Alternative OLS derivation

Suppose the population correlation between \mathbf{x} and ϵ is zero.

$$\begin{aligned}0 &= \mathbb{E} [\mathbf{x}\epsilon] \\0 &= \mathbb{E} [\mathbf{x} (y - \mathbf{x}'\beta)]\end{aligned}$$

Thus,

$$\beta = (\mathbb{E} [\mathbf{x}\mathbf{x}'])^{-1} \mathbb{E} [\mathbf{x}y]$$

If regression includes a constant, then these terms are covariance matrices, and we can use sample estimators in place of the population moments to get the OLS estimator:

$$\begin{aligned}\mathbf{b} &= (\mathbf{X}'\mathbf{X})^{-1} \mathbf{X}'\mathbf{Y} \\&= \left(\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i' \right)^{-1} \left(\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i y_i \right)\end{aligned}$$



Regression with an intercept

The assumption that X includes a column of 1's is important.

- ▶ Including a constant in the regression is equivalent to running a regression with demeaned data.
- ▶ Running a regression on just a constant regressor and nothing else, would simply pick up the mean in the data.
- ▶ Including a constant in the regression means the regressors try to match the variation in the y data, not the overall level.



Example: Risk premia

A fundamental theorem of asset pricing says that there is a linear relation between the risk premium of asset i , π_i , and a certain risk measure, x_i :

$$\pi_i = \alpha + \beta x_i + \epsilon_i$$

The Portfolio Theory class covers this theory in detail, but for now take it as given.

- ▶ Test this theory with a linear regression.
- ▶ Try both including a constant, α , and without.
- ▶ Risk and return data is collected on various industry portfolios.



Example: Regression with and without an intercept




Figure: Data Source: Ken French. Monthly 1926-2011.



THE UNIVERSITY OF
CHICAGO

Example: Regression with and without an intercept

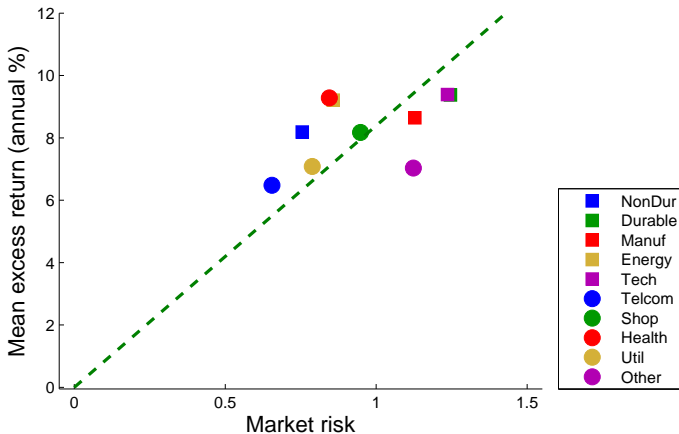


Figure: Data Source: Ken French. Monthly 1926-2011.



THE UNIVERSITY OF
CHICAGO

Example: Regression with and without an intercept




Figure: Data Source: Ken French. Monthly 1926-2011.



THE UNIVERSITY OF
CHICAGO

Example: Regression with and without an intercept




Figure: Data Source: Ken French. Monthly 1926-2011.



THE UNIVERSITY OF
CHICAGO

Residuals with zero mean

By assuming the model includes a constant,

$$\mathbb{E}[\mathbf{x}\epsilon] = \mathbf{0} \implies \mathbb{E}[\epsilon] = 0$$

By including a constant in the sample estimation,

$$\frac{1}{n} \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix}' \mathbf{e} = 0 \implies \frac{1}{n} \sum_{i=1}^n e_i = \bar{\mathbf{e}} = 0$$



R-squared

The **R-squared**, or coefficient of determination, in a regression is defined as

$$\begin{aligned} R_{y,x}^2 &= \frac{\text{regression sum of squares}}{\text{total sum of squares}} \\ &= 1 - \frac{\text{error sum of squares}}{\text{total sum of squares}} \end{aligned}$$

Algebraically, this is

$$\begin{aligned} R_{y,x}^2 &= \frac{\mathbf{b} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \\ &= 1 - \frac{\mathbf{e}'\mathbf{e}}{\sum_{i=1}^n (y_i - \bar{y})^2} \end{aligned}$$



R-squared versus correlation

Intuitively, the R-squared is the square of the correlation between y and the projection of y onto \mathbf{x} .

$$R_{y,\mathbf{x}}^2 = [\text{corr}(\mathbf{Y}, \mathbf{PY})]^2$$

In a univariate regression of y on x ,

$$R_{y,x}^2 = [\text{corr}(y, x)]^2$$



Caveat: Regressing on a constant

The interpretation and formula for R-squared does not hold if there is no constant regressor.

- ▶ Without a constant, the R-squared will not necessarily be between 0 and 1.
- ▶ Without a constant, the R-squared will not necessarily be the square of the correlation between the sample \mathbf{Y} and the projected \hat{Y} values.
- ▶ Without a regressor, the fit can be improved simply by shifting the sample \mathbf{Y} data by a constant.

